

# CSC570 NoSQL Databases

## Lab 1, Part 2: Relational Database Using PostgreSQL

### Part 2: Advanced Queries, Code, Rules

The tasks we want to complete in this section are to:

- Group similar values,
- Execute code on the server (stored procedures)
- Create custom interfaces using views and rules.

These are some of the features that make Relational Database so flexible in the ways it can query data. Stored procedures can help to move some business intelligence from the main application to the database- the idea is to have enough rules in the database that multiple applications can connect to it safely . Rules, and the ability to re-write them, enable other applications, some of which might be incompatible with the main business application, to use the business data. The database designer's task is to balance control and openness (which one depends on the situation).

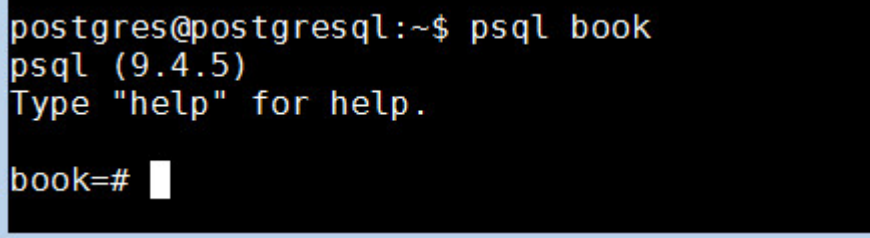
To get back to where we were at the end of part 1 of the lab:

1. Start your Back-end and Front-end VMs
2. On your FE machine, Double click on the "Turnkey PostgreSQL" Firefox Shortcut on your desktop.
3. Select the "Web Shell" interface. The login for your postgres user is:

Postgresql login: postgres

Password: p@ssw0rd

4. To drop into the Postgres command-line shell and connect to our created DB enter: `psql book`



```
postgres@postgresql:~$ psql book
psql (9.4.5)
Type "help" for help.

book=#
```

All of your prior work should still be present. You can quickly check by seeing if you country table is present by querying the book database with:

```
SELECT * FROM countries;
```

### Aggregate Functions

An aggregate query groups results from several rows by some common criteria. It can be as simple as counting the number of rows in a table or calculating the average of some numerical column. In order to explore this feature, we need insert some more data into our database.

Let's enter Sangamon Auditorium as a venue along with several upcoming events.

- a. Insert Springfield into the *cities* table and Sangamon Auditorium into the *venues* table:

```
book=# INSERT INTO cities VALUES ('Springfield', '62703', 'us');
INSERT 0 1
book=#
```

```
book=# INSERT INTO venues (name, postal_code, country_code)
book=# VALUES ('Sangamon Auditorium', '62703', 'us');
```

- b. Add the events to the *events* table:

(Here's some SQL that will set your venue\_id based on the name rather than trying to remember the number:

```
book=# INSERT INTO events(title,starts,ends,venue_id)
book=# VALUES ('Coldplay Concert', '2015-12-31 21:00:00', '2015-12-31 23:30:00', (
book=# SELECT venue_id FROM venues WHERE name ='Crystal Ballroom'));
```

Your event table should have the following:

event_id	title	starts	ends	venue_id
1	LARP Club	2012-02-15 17:30:00	2012-02-15 19:30:00	2
3	Christmas Day	2012-12-25 00:00:00	2012-12-25 23:59:00	
4	Coldplay Concert	2015-12-31 21:00:00	2015-12-31 23:30:00	1
5	Annie	2016-02-13 15:00:00	2016-02-13 18:30:00	3
6	42nd Street	2016-03-06 19:00:00	2016-02-13 22:30:00	3
7	The Chieftains	2016-03-01 19:00:00	2016-02-13 22:30:00	3
2	April Fools Day	2012-04-01 00:00:00	2012-04-01 23:59:00	

(7 rows)

### Aggregate Function: COUNT()

The simplest aggregate function is count(), which returns the integer count of the occurrences of something . Counting the total number of events in our events table should return 7:

```
book=# SELECT count(*) FROM events;
count
-----
      7
(1 row)
```

A more useful count would be all event titles that contain the string "Day" (note: % is a wildcard on LIKE searches), you should receive a value of 2.

```
book=# SELECT count(title) FROM events WHERE title LIKE '%Day%';
count
-----
      2
(1 row)
```

Or, to get the first start time and last end time of all events at the Sangamon Auditorium, use min() (return the smallest value) and max() (return the largest value):

```
book=# SELECT min(starts), max(ends) FROM events INNER JOIN venues
book-# ON events.venue_id=venues.venue_id
book-# WHERE venues.name='Sangamon Auditorium';
      min      |      max
-----+-----
2016-02-13 15:00:00 | 2016-02-13 22:30:00
(1 row)
```

Aggregate functions are useful but limited on their own. If we wanted to count all events at each venue, we could write the following for each venue ID:

```
SELECT count(*) FROM events WHERE venue_id = 1;
SELECT count(*) FROM events WHERE venue_id = 2;
SELECT count(*) FROM events WHERE venue_id = 3;
SELECT count(*) FROM events WHERE venue_id IS NULL;
```

This would be tedious and would become unmanageable as the number of venues grows. Thankfully, there is the GROUP BY aggregate.

### *Aggregate Function: GROUP BY ()*

GROUP BY is a shortcut for running the previous queries all at once. With GROUP BY, you tell Postgres to place the rows into groups and then perform some aggregate function (such as count()) on those groups.

```
book=# SELECT venue_id, count(*) FROM events GROUP BY venue_id;
 venue_id | count
-----+-----
        1 |      2
        3 |      1
        2 |      3
        2 |      1
(4 rows)
```

GROUP BY can be made more flexible by using a filter. The GROUP BY condition has its own filter keyword: HAVING. HAVING is like the WHERE clause, except it can filter by aggregate functions (whereas WHERE cannot). The following query SELECTs the most popular venues, those with two or more events:

```
book=# SELECT venue_id FROM events GROUP BY venue_id
book-# HAVING count(*) >=2 AND venue_id IS NOT NULL;
 venue_id
-----
        3
(1 row)
```

You can use GROUP BY without any aggregate functions. If you call SELECT...FROM...GROUP BY on one column, you get all unique values.

```
SELECT venue_id FROM events GROUP BY venue_id;
```

This kind of grouping is so common that SQL has a shortcut in the `DISTINCT` keyword.

```
SELECT DISTINCT venue_id FROM events;
```

The results of both queries will be identical.

## Transactions

Transactions are the bulwark of relational database consistency. Transactions ensure that every command of a set is executed. If anything fails along the way, all of the commands are rolled back like they never happened. PostgreSQL transactions follow ACID compliance, which stands for Atomic (all ops succeed or none do), Consistent (the data will always be in a good state—no inconsistent states), Isolated (transactions don't interfere with each other), and Durable (a committed transaction is safe, even after a server crash).

We can wrap any transaction within a `BEGIN TRANSACTION` block. To verify atomicity, we end the transaction with the `ROLLBACK` command.

```
BEGIN TRANSACTION;
DELETE FROM events;
ROLLBACK;
```

```
book=# BEGIN TRANSACTION;
BEGIN
book=# DELETE FROM events;
DELETE 7
book=# ROLLBACK;
ROLLBACK
```

The events all remain. If we run `SELECT * FROM events;` all of our data is still there

```
book=# SELECT * FROM events;
 event_id | title           | starts           | ends           | venue_id
-----+-----+-----+-----+-----
      1 | LARP Club       | 2012-02-15 17:30:00 | 2012-02-15 19:30:00 |      2
      3 | Christmas Day   | 2012-12-25 00:00:00 | 2012-12-25 23:59:00 |
      4 | Coldplay Concert | 2015-12-31 21:00:00 | 2015-12-31 23:30:00 |      1
      5 | Annie           | 2016-02-13 15:00:00 | 2016-02-13 18:30:00 |      3
      6 | 42nd Street     | 2016-03-06 19:00:00 | 2016-02-13 22:30:00 |      3
      7 | The Chieftains  | 2016-03-01 19:00:00 | 2016-02-13 22:30:00 |      3
      2 | April Fools Day | 2012-04-01 00:00:00 | 2012-04-01 23:59:00 |
(7 rows)
```

Transactions are useful when you're modifying two tables that you don't want out of sync. The classic example is a debit/credit system for a bank, where money is moved from one account to another:

```
BEGIN TRANSACTION;
UPDATE account SET total=total+5000.0 WHERE account_id=1337;
UPDATE account SET total=total-5000.0 WHERE account_id=45887;
END;
```

If something happened between the two updates, this bank would lose \$5000.00. But, when wrapped in a transaction block, the initial update is rolled back, even if the server crashes.

## Stored Procedures and Triggers

All of the commands we've used so far have been declarative, but sometimes you may need to run some code. As a Database Designer, you would now have to make a decision: execute code on the client side or execute code on the database side (server side).

Stored procedures can offer large performance advantages but at a large architectural costs. You may avoid streaming thousands of rows to a client application, but now you have bound your application code to this database. The decision whether or not to stored procedures is an important one.

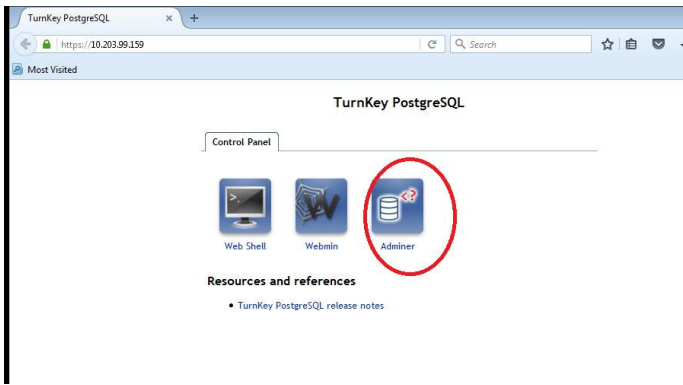
#### Why use Stored Procedures?

- i. Reusability
  - a. Avoid rewriting subqueries and improve readability.
  - b. If you can't store a query in a library that all the applications can access, you can put that query in a stored procedure.
- ii. Separation of duties
  - a. You don't trust non-DBA's to write queries.
- iii. Data integrity
  - a. Use triggers or constraints to prevent bad data from entering.
  - b. Run several interdependent queries in a transaction in a single stored procedure.
- iv. Event handling
  - a. Log changes.
  - b. Notify other systems of new data.

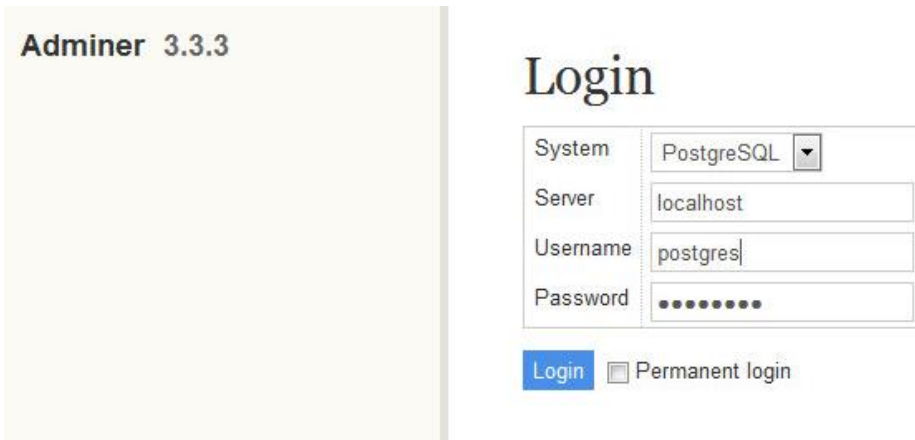
#### Why NOT use Stored Procedures?

- 1. Views may be all you need.
  - a. An object-relational mapper (ORM) can help write queries safely.
  - b. Difficult to version control stored procedures.
  - c. Software rollouts may require more db changes.
  - d. Could slow software development process.

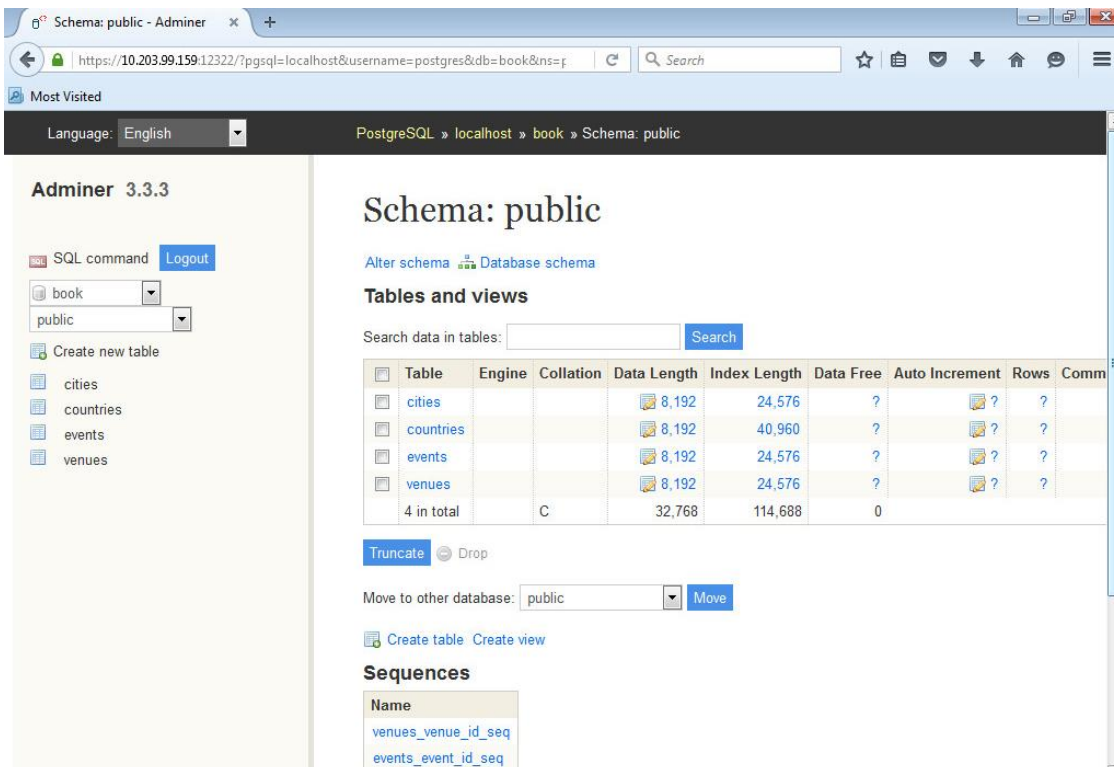
To look at the power of using stored procedures, let's use a Database Admin tool. Adminer has been installed on your machine and is available from your Mozilla shortcut:



You login with Username: postgres and Password: p@ssw0rd



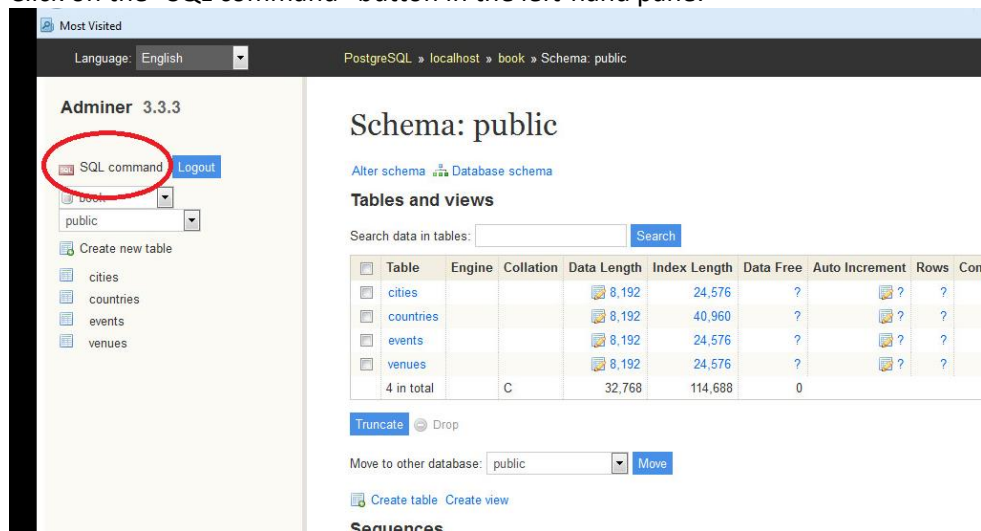
You should now see all of the databases installed on our server. We want to select our “book” database; double-click on it to see our existing tables:





Admin tools, like Adminer, as typically used by DBAs to manage databases. We could have done all of our Database CRUD from a tool, but each tool is slightly different. Working from a command line is constant for all installations.

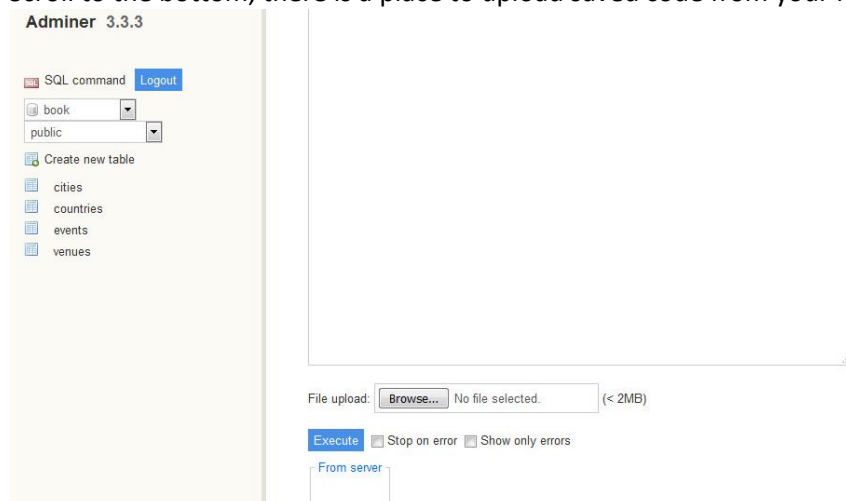
Click on the “SQL command” button in the left-hand pane:



This will open up a textbox in which we can enter SQL code:



Scroll to the bottom, there is a place to upload saved code from your FE machine.



Now let's create a stored procedure (or FUNCTION) that simplifies INSERTing a new event at a venue without needing the venue\_id. Further, if the user tries to enter an event for a venue that doesn't exist, the procedure will create it first and reference it in the new event. Also, we'll return a Boolean indicating whether a new venue was added, as a check to let the user know a new venue has been created (and correct the input if that's not what they wanted).

```
CREATE OR REPLACE FUNCTION add_event( title text, starts timestamp,
ends timestamp, venue text, postal varchar(9), country char(2) )
RETURNS boolean AS $$
DECLARE
    did_insert boolean := false;
    found_count integer;
    the_venue_id integer;
BEGIN
    SELECT venue_id INTO the_venue_id
    FROM venues v
    WHERE v.postal_code=postal AND v.country_code=country AND v.name ILIKE venue
    LIMIT 1;

    IF the_venue_id IS NULL THEN
        INSERT INTO venues (name, postal_code, country_code)
        VALUES (venue, postal, country)
        RETURNING venue_id INTO the_venue_id;

        did_insert := true;
    END IF;

    RAISE NOTICE 'Venue found %', the_venue_id;

    INSERT INTO events (title, starts, ends, venue_id)
    VALUES (title, starts, ends, the_venue_id);

    RETURN did_insert;
END;
$$ LANGUAGE plpgsql;
```

You can type it in the SQL command box in Adminer or download the .sql file from Blackboard onto your Front-end machine. You can then copy from the .sql file into the command window or upload it to Execute. (If you upload it, you will not see the code, it will only tell you that it successfully executed.)



## Adminer 3.3.3

SQL command Logout

book

public

Create new table

cities

countries

events

venues

1 query executed OK. (0.016 s)

```

CREATE OR REPLACE FUNCTION add_event( title text, starts timestamp,
ends timestamp, venue text, postal varchar(9), country char(2) )
RETURNS boolean AS $$
DECLARE
    did_insert boolean := false;
    found_count integer;
    the_venue_id integer;
BEGIN
    SELECT venue_id INTO the_venue_id
    FROM venues v
    WHERE v.postal_code=postal AND v.country_code=country AND v.name ILIKE venue
    LIMIT 1;

    IF the_venue_id IS NULL THEN
        INSERT INTO venues (name, postal_code, country_code)
        VALUES (venue, postal, country)
        RETURNING venue_id INTO the_venue_id;

        did_insert := true;
    END IF;

    -- Note: not an "error", as in some programming languages
    RAISE NOTICE 'Venue found %', the_venue_id;

    INSERT INTO events (title, starts, ends, venue_id)
    VALUES (title, starts, ends, the_venue_id);

    RETURN did_insert;
END;
$$ LANGUAGE plpgsql;

```

How can we check if our stored procedure works? Let's use it to add a new event in a new venue: UIS' May Graduation at the Prairie Capital Convention Center:

```
SELECT add_event('UIS ,, Graduation', '2016-05-14 16:00:00', '2016-05-14 20:00:00', 'Prairie Capital Convention Center', '62703', 'us');
```

Adminer executes the saved add\_event procedure:

PostgreSQL » localhost » book » public » SQL command

## SQL command

```
SELECT add_event('UIS Graduation', '2016-05-14 16:00:00', '2016-05-14 20:00:00', 'Prairie Capital Convention Center', '62703', 'us');
```

add\_event

t

1 row (0.009 s) Edit, EXPLAIN, Export

```
SELECT add_event('UIS Graduation', '2016-05-14 16:00:00', '2016-05-14 20:00:00', 'Prairie Capital Convention Center', '62703', 'us');
```

We can examine our tables in Adminer to see that the new event and the new venue have been added:

**Adminer 3.3.3**

SQL command [Logout](#)

book public

Create new table

- cities
- countries
- events
- venues

### Select: events

Select data Show structure Alter table New item

Select Search Sort Limit 30 Text length 100 Action Select

```
>> SELECT * FROM "events" LIMIT 30 Edit
```

edit	event_id	title	starts	ends	venue_id
	1	LARP Club	2012-02-15 17:30:00	2012-02-15 19:30:00	2
	3	Christmas Day	2012-12-25 00:00:00	2012-12-25 23:59:00	NULL
	4	Coldplay Concert	2015-12-31 21:00:00	2015-12-31 23:30:00	1
	5	Annie	2016-02-13 15:00:00	2016-02-13 18:30:00	3
	6	42nd Street	2016-03-06 19:00:00	2016-02-13 22:30:00	3
	7	The Chieftains	2016-03-01 19:00:00	2016-02-13 22:30:00	3
	2	April Fools Day	2012-04-01 00:00:00	2012-04-01 23:59:00	NULL
	9	UIS Graduation	2016-05-14 16:00:00	2016-05-14 20:00:00	

(8 rows) whole result

Edit Save Edit Clone Delete Export Import

---

**Adminer 3.3.3**

SQL command [Logout](#)

book public

Create new table

- cities
- countries
- events
- venues

### Select: venues

Select data Show structure Alter table New item

Select Search Sort Limit 30 Text length 100 Action Select

```
>> SELECT * FROM "venues" LIMIT 30 Edit
```

edit	venue_id	name	street_address	type	postal_code	country_code
	1	Crystal Ballroom	NULL	public	97205	us
	2	Voodoo Donuts	NULL	public	97205	us
	3	Sangamon Auditorium	NULL	public	62703	us
	4	Pratt Capital Convention Center	NULL	public	62705	us

(4 rows) whole result

Edit Save Edit Clone Delete Export Import

The power of a stored procedure is when they are combined with triggers. Triggers automatically launch stored procedures when some event happens, like an insert or update. They allow the database to automatically enforce some required behavior in response to changing data.

Let's create a new stored procedure that logs whenever an event is updated (we want a record of anytime someone alters an existing event). First, create a logs table to store event changes. A primary key isn't necessary here, since it's just a log. Click on the "Create new table" button in the left-hand pane and enter the following in the textboxes:

**Adminer 3.3.3**

SQL command [Logout](#)

book public

Create new table

- cities
- countries
- events
- venues

### Create table

Table name: logs Save

Column name	Type	Length	Options	NULL	AI	Default values	+
event_id	integer			<input type="checkbox"/>	<input type="radio"/>		x
old_title	character varying	255		<input type="checkbox"/>	<input type="radio"/>		x
old_starts	timestamp			<input type="checkbox"/>	<input type="radio"/>		x
old_ends	timestamp			<input type="checkbox"/>	<input type="radio"/>		x
logged_at	timestamp			<input type="checkbox"/>	<input type="radio"/>	current_timestamp	x

Auto Increment: ☐ Default values ☒ Comment ☐

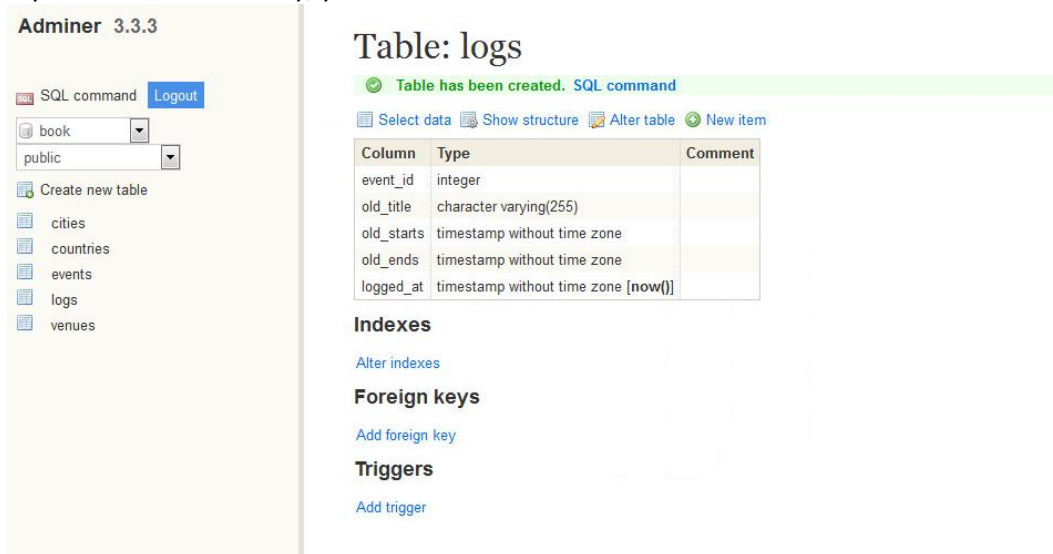
Save

If you are having trouble reading the textboxes above, I have entered the options below:

```
CREATE TABLE logs (  
  event_id integer,  
  old_title varchar(255),  
  old_starts timestamp,  
  old_ends timestamp,  
  logged_at timestamp DEFAULT current_timestamp  
);
```

You could also enter the above code in the SQL command window and it will create the table.

If you created it correctly, you should see:



The screenshot shows the Adminer 3.3.3 web interface. On the left is a sidebar with a 'SQL command' tab and a 'Logout' button. Below these are dropdown menus for 'book' and 'public', and a 'Create new table' button. A list of existing tables (cities, countries, events, logs, venues) is shown. The main panel is titled 'Table: logs' and displays a success message: 'Table has been created. SQL command'. Below this are links for 'Select data', 'Show structure', 'Alter table', and 'New item'. A table structure is shown with columns: event\_id (integer), old\_title (character varying(255)), old\_starts (timestamp without time zone), old\_ends (timestamp without time zone), and logged\_at (timestamp without time zone [now()]). Below the table structure are sections for 'Indexes', 'Foreign keys', and 'Triggers', each with an 'Add' link.

Column	Type	Comment
event_id	integer	
old_title	character varying(255)	
old_starts	timestamp without time zone	
old_ends	timestamp without time zone	
logged_at	timestamp without time zone [now()]	

Next, we build a function to insert old data into the log. The OLD variable represents the row about to be changed (NEW represents an incoming row. It will then Output a notice to the console with the event\_id before returning.

```
CREATE OR REPLACE FUNCTION log_event() RETURNS trigger AS $$  
DECLARE  
BEGIN  
  INSERT INTO logs (event_id, old_title, old_starts, old_ends)  
  VALUES (OLD.event_id, OLD.title, OLD.starts, OLD.ends);  
  RAISE NOTICE 'Someone just changed event #%', OLD.event_id;  
  RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;
```

Type this into your SQL command window and execute:

Language: English
PostgreSQL » localhost » book » public » SQL command

**Adminer 3.3.3**

SQL command
Logout

book
public

Create new table

cities
countries
events
logs

## SQL command

```

CREATE OR REPLACE FUNCTION log_event() RETURNS trigger AS $$
DECLARE
BEGIN
INSERT INTO logs (event_id, old_title, old_starts, old_ends)
VALUES (OLD.event_id, OLD.title, OLD.starts, OLD.ends);
RAISE NOTICE 'Someone just changed event #%', OLD.event_id;
RETURN NEW;
END;
$$ LANGUAGE plpgsql;

```

Finally, we create our trigger to log changes after any row is updated by entering and executing the following SQL:

```

CREATE TRIGGER log_events
AFTER UPDATE ON events
FOR EACH ROW EXECUTE PROCEDURE log_event();

```

Language: English
PostgreSQL » localhost » book » public » SQL command

**Adminer 3.3.3**

SQL command
Logout

book
public

Create new table

cities
countries
events
logs
venues

## SQL command

```

CREATE TRIGGER log_events AFTER UPDATE ON events
FOR EACH ROW EXECUTE PROCEDURE log_event();

```

Now let's test it by updating the time Graduation will end:

```

UPDATE events
SET ends='2016-05-14 19:00:00'
WHERE title='UIS Graduation';

```

## SQL command

```

UPDATE events SET ends='2016-05-14 19:00:00' WHERE title='UIS Graduation'

```

Query executed OK, 1 row affected, (0.006 s)
Edit

```

UPDATE events SET ends='2016-05-14 19:00:00' WHERE title='UIS Graduation';

```

We can see this change in our *logs* table:

# Select: logs

[Select data](#) [Show structure](#) [Alter table](#) [New item](#)

Select	Search	Sort	Limit 30	Action <a href="#">Select</a>
--------	--------	------	-------------	----------------------------------

>> `SELECT * FROM "logs" LIMIT 30` [Edit](#)

<a href="#">edit</a>	event_id	old_title	old_starts	old_ends	logged_at
<a href="#">edit</a>	9	UIS Graduation	2016-05-14 16:00:00	2016-05-14 20:00:00	2016-02-06 23:22:59.732369

(1 row) [whole result](#)

Edit	<a href="#">Export</a>	<a href="#">Import</a>
<a href="#">Save</a> <a href="#">Edit</a> <a href="#">Clone</a> <a href="#">Delete</a>		

Triggers can also be created before updates and before or after inserts. Triggers are tied to the table whose Update or Insert causes its execution. In our example it is a change on *events*. Adminer stores the information for the trigger with its triggering table:

Language: English

PostgreSQL » localhost » book » public » Table: events

Adminer 3.3.3

[SQL command](#) [Logout](#)

book

public

Create new table

cities

countries

events

logs

venues

Table: events

[Select data](#) [Show structure](#) [Alter table](#) [New item](#)

Column	Type	Comment
event_id	integer Auto Increment [nextval('events_event_id_seq')]	
title	text NULL	
starts	timestamp without time zone NULL	
ends	timestamp without time zone NULL	
venue_id	integer NULL	

Indexes

PRIMARY	event_id
---------	----------

[Alter indexes](#)

Foreign keys

Source	Target	ON DELETE	ON UPDATE
venue_id	venues (venue_id)		<a href="#">Alter</a>

[Add foreign key](#)

Triggers

UPDATE	log_events	<a href="#">Alter</a>
--------	------------	-----------------------

[Add trigger](#)

## Views and Rules

Relational Database allows for the creation of very complex queries. The results are often a subset or union of data that we would like to treat just like any other table. VIEWS allow us to do that. Unlike stored procedures, these aren't functions being executed but rather aliased queries. In our current database, all holidays contain the word Day and have no venue. Creating a view on just holidays is as simple as writing a query and prefixing it with CREATE VIEW *some\_view\_name* AS...

```
CREATE VIEW holidays AS
SELECT event_id AS holiday_id, title AS name, starts AS date
FROM events
WHERE title LIKE '%Day%' AND venue_id IS NULL;
```

The screenshot shows the Adminer 3.3.3 interface. On the left, the 'SQL command' tab is active, and the 'book' database is selected. The 'public' schema is also selected. The 'holidays' table is highlighted in the left sidebar. The main window displays the SQL command: `CREATE VIEW holidays AS SELECT event_id AS holiday_id, title AS name, starts AS date FROM events WHERE title LIKE '%Day%' AND venue_id IS NULL;`. A green message bar indicates 'Query executed OK, 0 rows affected. (0.031 s)'. Below the command, the same SQL command is repeated in a text area.

Let's test it by querying the new holiday view:

```
SELECT name, to_char(date, 'Month DD, YYYY') AS date
FROM holidays
WHERE date <= '2012-12-25';
```

The screenshot shows the Adminer 3.3.3 interface. The 'SQL command' tab is active, and the 'book' database is selected. The 'public' schema is also selected. The 'holidays' table is highlighted in the left sidebar. The main window displays the SQL command: `SELECT name, to_char(date, 'Month DD, YYYY') AS date FROM holidays WHERE date <= '2012-12-25';`. The query results are displayed in a table with two columns: 'name' and 'date'. The results show 'Christmas Day' and 'December 25, 2012'. A green message bar indicates '1 row (0.006 s)'. Below the command, the same SQL command is repeated in a text area.

Views are powerful tools for opening up complex queried data in a simple way. If you want to add a new column to the view, it will have to come from the underlying table. Let's alter the events table to have an array of associated colors.



ALTER TABLE events  
ADD colors text ARRAY;



Since holidays are to have colors associated with them, let's update the VIEW query to contain the colors array.

```
CREATE OR REPLACE VIEW holidays AS  
SELECT event_id AS holiday_id, title AS name, starts AS date, colors  
FROM events  
WHERE title LIKE '%Day%' AND venue_id IS NULL;
```



Now we can add color strings to the holiday of choice.

Unfortunately, we cannot update a view directly. For this we need a RULE. A RULE is a description of how to alter the parsed query tree. Every time Postgres runs an SQL statement, it parses the statement into a query tree (generally called an abstract syntax tree). So, to allow updates against our holidays view, we need to craft a RULE that tells Postgres what to do with an UPDATE. Our rule will capture updates to the holidays view and instead run the update on events, pulling



values from the pseudorelations NEW and OLD. NEW functionally acts as the relation containing the values we're setting, while OLD contains the values we query by.

```
CREATE RULE update_holidays AS ON UPDATE TO holidays DO INSTEAD
UPDATE events
SET title = NEW.name,
    starts = NEW.date,
    colors = NEW.colors
WHERE title = OLD.name;
```

The screenshot shows the Adminer 3.3.3 interface. On the left is a sidebar with a 'SQL command' tab and a list of tables: cities, countries, events, holidays, logs, and venues. The main panel is titled 'SQL command' and contains the following SQL code:

```
CREATE RULE update_holidays AS ON UPDATE TO holidays DO INSTEAD
UPDATE events
SET title = NEW.name,
    starts = NEW.date,
    colors = NEW.colors
WHERE title = OLD.name;
```

Below the code, a green status bar indicates: 'Query executed OK, 0 rows affected. (0.007 s)'. A second, identical SQL command is shown in a preview area below.

With this rule in place, now we can update holidays directly.  
UPDATE holidays SET colors = '{"red","green"}' where name = 'Christmas Day';

The screenshot shows the Adminer 3.3.3 interface. The main panel is titled 'SQL command' and contains the following SQL code:

```
UPDATE holidays SET colors= '{"red","green"}' WHERE name='Christmas Day';
```

Below the code, a green status bar indicates: 'Query executed OK, 1 row affected. (0.017 s)'. A second, identical SQL command is shown in a preview area below.

Viewing the data in our tables, we can see it worked on both *holidays* and *events*:

The image contains two side-by-side screenshots of the Adminer 3.3.3 interface showing data in the *holidays* and *events* tables.

**Select: holidays**

The screenshot shows the 'Select: holidays' view. The table has columns: edit, holiday\_id, name, date, and colors. The data is as follows:

edit	holiday_id	name	date	colors
	2	April Fools Day	2012-04-01 00:00:00	NULL
	3	Christmas Day	2012-12-25 00:00:00	{red,green}

The row for Christmas Day (id 3) is circled in red. The status bar shows '(2 rows) whole result'.

**Select: events**

The screenshot shows the 'Select: events' view. The table has columns: edit, event\_id, title, starts, ends, venue\_id, and colors. The data is as follows:

edit	event_id	title	starts	ends	venue_id	colors
	1	LARP Club	2012-02-15 17:30:00	2012-02-15 19:30:00	2	NULL
	4	Coldplay Concert	2015-12-31 21:00:00	2015-12-31 23:30:00	1	NULL
	5	Annie	2016-02-13 15:00:00	2016-02-13 18:30:00	3	NULL
	6	42nd Street	2016-03-06 19:00:00	2016-02-13 22:30:00	3	NULL
	7	The Chieftains	2016-03-01 19:00:00	2016-02-13 22:30:00	3	NULL
	2	April Fools Day	2012-04-01 00:00:00	2012-04-01 23:59:00	NULL	NULL
	9	UIS Graduation	2016-05-14 16:00:00	2016-05-14 19:00:00	6	NULL
	3	Christmas Day	2012-12-25 00:00:00	2012-12-25 23:59:00	NULL	{red,green}

The row for Christmas Day (id 3) is circled in red. The status bar shows '(8 rows) whole result'.

## *Work on your own to submit*

This is a quick lesson on both relational database and the tools that are used by database admins to work with them. The following are some final exercises to let you try these advanced features on your own.

### 1. Views

To submit your work, add this code at the bottom of the file for Week 1 (Lab 1 Part 1). The file should be named lab1.sql and uploaded to GitHub.

Your code should be lightly commented. Use `--` to add comments (double dashes) that explain each step listed below. This is good practice for when you share code with colleagues so that they can contribute to the project.

- a. Using the tables created in Lab 1 Part 1, insert 20 transactions. Three of these transactions need to have the `actual_checkin_time` after the `scheduled_checkin_time`. This will allow you to test the view you will be creating in the next steps. For example, a transaction where the `scheduled_checkin_time` is 2018-08-01 14:39:53 and the `actual_checkin_time` is 2018-08-02 14:39:53. Additionally, five of the transactions need to have a `checkout_time` after September 3 2018.
- b. Create a late checkins view of distinct items that were checked in late grouped by `user_id`, `inventory_id`, and description. This view should display the total number of late checkins per device per user. For example, if user1 checked in two items late, there should be two rows displayed for user1 and each row should include the total number of times that user returned that particular item late.
- c. Test the late checkins view by selecting and displaying all records from the view.

### 2. The tasks we want to complete in the second part of Lab 1 Part 2 are to:

- a. On our PostgreSQL VM, run the lab1.sql file from Step 1 to create the tables and data.
- b. Access our PostgreSQL databases from a Laravel PHP app. Laravel is a PHP framework.
- c. Execute a query from Laravel using Laravel's Eloquent Object-Relational Mapping (ORM).
- d. Display our query results on a webpage.

Utilizing a database might be all that you need in the real world, but accessing that database from a programming language such as PHP will give you an advantage in the workplace. Knowing how to retrieve, manipulate and save data from a programming language is not a new feature of any programming language or database, so it will be helpful to understand how to set-up and perform some basic tasks, such as the ones we will cover now:

1. Install Visual Studio code integrated development environment (IDE). You can use other PHP IDE's such as Atom.
2. Map a drive to the Windows share on the container with the Laravel code
3. Open up the Laravel project in the Visual Studio code IDE
4. Modify the Laravel controller and view to execute the queries using the ORM relationships and display the data in the lab

## Video Tutorials

Please review these video lectures:

### Lab Setup

[9 min 33 sec]:

[https://cdnapisec.kaltura.com/index.php/extwidget/preview/partner\\_id/1371761/uiconf\\_id/31473632/entry\\_id/1\\_n4lu2j8q/embed/dynamic](https://cdnapisec.kaltura.com/index.php/extwidget/preview/partner_id/1371761/uiconf_id/31473632/entry_id/1_n4lu2j8q/embed/dynamic)

### Access from Programming Language (PHP/Laravel Eloquent ORM)

[42 min 31 sec]:

[https://cdnapisec.kaltura.com/index.php/extwidget/preview/partner\\_id/1371761/uiconf\\_id/31473632/entry\\_id/1\\_yx21kgc5/embed/dynamic](https://cdnapisec.kaltura.com/index.php/extwidget/preview/partner_id/1371761/uiconf_id/31473632/entry_id/1_yx21kgc5/embed/dynamic)

Please be sure to follow the steps in the Lab Setup video linked above to connect to your container. Use this table to determine which container is yours. You will log into the share with `.\\NetID` for the username (`.\\tllos1` for example) and **your UIN for the password**.

<u>Netid</u>	<u>Windows share</u>	<u>Url of the PHP application</u>
agang2	\\10.64.3.56\\agang2	<a href="https://csc570e.uis.edu:9444">https://csc570e.uis.edu:9444</a>
bbala5	\\10.64.3.56\\bbala5	<a href="https://csc570e.uis.edu:9445">https://csc570e.uis.edu:9445</a>
bguti6	\\10.64.3.56\\bguti6	<a href="https://csc570e.uis.edu:9446">https://csc570e.uis.edu:9446</a>
brodr22	\\10.64.3.56\\brodr22	<a href="https://csc570e.uis.edu:9447">https://csc570e.uis.edu:9447</a>
chick7	\\10.64.3.56\\chick7	<a href="https://csc570e.uis.edu:9448">https://csc570e.uis.edu:9448</a>
eunsik2	\\10.64.3.56\\eunsik2	<a href="https://csc570e.uis.edu:9449">https://csc570e.uis.edu:9449</a>
jlund6	\\10.64.3.56\\jlund6	<a href="https://csc570e.uis.edu:9450">https://csc570e.uis.edu:9450</a>
jshei3	\\10.64.3.56\\jshei3	<a href="https://csc570e.uis.edu:9451">https://csc570e.uis.edu:9451</a>
mpavl3	\\10.64.3.56\\mpavl3	<a href="https://csc570e.uis.edu:9452">https://csc570e.uis.edu:9452</a>
rsayy2	\\10.64.3.56\\rsayy2	<a href="https://csc570e.uis.edu:9453">https://csc570e.uis.edu:9453</a>
sarya7	\\10.64.3.56\\sarya7	<a href="https://csc570e.uis.edu:9454">https://csc570e.uis.edu:9454</a>
skoch7	\\10.64.3.56\\skoch7	<a href="https://csc570e.uis.edu:9455">https://csc570e.uis.edu:9455</a>
szhen6	\\10.64.3.56\\szhen6	<a href="https://csc570e.uis.edu:9456">https://csc570e.uis.edu:9456</a>
zwold2	\\10.64.3.56\\zwold2	<a href="https://csc570e.uis.edu:9457">https://csc570e.uis.edu:9457</a>
smeka6	\\10.64.3.56\\smeka6	<a href="https://csc570e.uis.edu:9458">https://csc570e.uis.edu:9458</a>

Relationships will need to be set before writing the queries (`hasMany`, `belongsTo`, etc). For Step 1, you will need to create a relationship from the Users table to the Transactions table, from the Transactions table to the Users table, and from the Transactions table to the Inventory table. For Step 2, you will need to create a relationship from the Transactions table to the User table and from the Transactions table to the Inventory table.

Your results will be the view `transaction.index`, which already maps to the route of the root of the website. This view will display two tables with the data similar to what I showed in the video example for Status and Inventory (see video link above for Access from Programming Language).

1. Write a query to show all items that user1 has checked out with the corresponding `checkout_time` and display the results.
2. Write a query to show all users who checked out items before September 3 2018 and display the results, which should include the user first name, the item description, the `checkout_time`, and the status.

